

Fall 2018 Comp 430/533
Assignment 6
Due: Date of course final.

A6: NoSQL

The goal of this assignment is to explore using a NoSQL database.

Grading

This assignment is worth 100 points.

Be sure to include both your code and your output in your submission.

Academic Honesty

The following level of collaboration is allowed on this assignment:

You may discuss the assignment with your classmates at a high level. Any issues getting MongoDB running is totally fine. What is not allowed is direct examination of anyone else's NoSQL code (on a computer, email, whiteboard, etc.) or allowing anyone else to see your NoSQL code. You may also use (and in fact are encouraged to use) the MongoDB reference manual

<https://docs.mongodb.com/manual/>

You may use the search engine of your choice to look up the syntax for MongoDB commands, but may not use it to find answers.

You MAY post and discuss results with your classmates.

Loading the data

1. Start your mongodb server

```
mongod -auth -port 27017 -dbpath /data/db
```

2. Load the data files into MongoDB using the mongoimport program

- (a) `mongoimport --db ricedb --collection foodInfo --jsonArray --file foodInfo.json --username ricedb --password myPwd`
2012 documents.
- (b) `mongoimport --db ricedb --collection truckEvent --jsonArray --file truckEvent.json --username ricedb --password myPwd`
63 documents.
- (c) `mongoimport --db ricedb --collection sale --jsonArray --file sale.json --username ricedb --password myPwd`
6807 documents.
- (d) `mongoimport --db ricedb --collection ingredient --jsonArray --file ingredient.json --username ricedb --password myPwd`
78 documents.
- (e) `mongoimport --db ricedb --collection unit --jsonArray --file unit.json --username ricedb --password noSqlIsBest`
86 documents.

If you need to reload the data, empty the collections by running the following commands in the mongodb client

```
db.foodInfo.drop()
db.truckEvent.drop()
db.sale.drop()
db.ingredient.drop()
db.unit.drop()
```

1 Queries

Be sure that you are using the ricedb database!

```
use ricedb
```

1. (2 points) Write a query that returns a single document from the food-Info collection.

2. (4 points) Write a query that pretty prints (using the "pretty()" function) all documents with `portion_display_name = "regular Oreo"`.
3. (4 points) Write a query that pretty prints all the documents whose `display_name` contains "ice cream" (case insensitive) and that have less than 200 calories.
4. (4 points) Write a query that pretty prints the 5 `truckEvent` documents that describe events that started closest in time to, but strictly before, September 7, 2017. Order the events from closest time to furthest away.

In other words, if we have:

Day1 Day2 Day3 Day4 Day5 Day6 Day7 Day8 Day9 Day10

and we ask for the 3 days before Day7, you would return:

Day6

Day5

Day4

5. (4 points) Write a query that returns the total number of products sold during the `truckEvent` with `eventId` 16.
6. (8 points, 2 points each sub part) There are 4 parts to this question.
 - (a) Change the name of a `truckEvent` document where the `eventName` is "Pi Day" to "Pie Day". Include your query here.
 - (b) Write a query that pretty prints all `truckEvent` documents where the `eventName` is "Pi Day" OR "Pie Day".
 - (c) Change the `eventName` of all `truckEvent` documents where the `eventName` is "GSA Coffee Break" to "GSA Study Break".
 How many documents were updated? (Paste your MongoDB command and the output of your MongoDB command)
 - (d) Write a query that returns the number of `truckEvent` documents with `eventName` "GSA Study Break" OR "GSA Coffee Break". Include the number of documents found.

7. We want to learn about the calories in each of the different toppings. To help us figure this out, we have two new collections: `foodInfo` and `unit`. The `foodInfo` collection contains nutritional information about commonly consumed items. The `unit` collection maps different portion names to common unit measures (e.g. cup to ounces).

We want to match up toppings in our ingredient collection with toppings in the `foodInfo` collection. However, there are some challenges. Since there are multiple `foodInfo` documents with the same `display_name`, we need to disambiguate. One approach is to create a new collection that only has documents where `category = "topping"` from the `foodInfo` collection.

- (a) (6 points) Using an aggregate function, create a new collection, called `"foodInfoToppings"` that contains only the `display_name`, `portion_display_name`, and `calories` fields from `foodInfo`. Include every document where `category = "topping"`. Include your aggregation code in your submission as well as a query to obtain the results.
- (b) (15 points) Now, we want to know which toppings have the highest calories per ounce. Using the aggregation function that we went over in class, filter the ingredient collection to only contain documents where `category = "topping"` and then "join" these results to the `foodInfoToppings` collection. Note that the `ingName` field in ingredient will match the `display_name` field in `foodInfoToppings`.

Next, "join" the resulting dataset to the `unit` collection to get the number of ounces for each topping. Compute the calories per ounce (use the `$divide` aggregation function to do the division) and extract the following fields:

- `ingId`
- `ingName`
- `portion_display_name`
- `calories`
- `ounces`
- `calPerOunce` (computed)

Sort by `calPerOunce`, descending, and pretty print the three toppings with the highest calorie per ounce.

2 Streaming Twitter Data

As we discussed in class, Document stores are designed to handle streaming data in formats that can vary. For the second part of this assignment, you will collect some twitter data about ice cream and / or food trucks, store it in MongoDB and then query the data.

We will use the tweepy Python library to collect tweets (<http://www.tweepy.org/>).

1. Create a twitter account, if you don't already have one. Go to <https://developer.twitter.com/content/developer-twitter/en.html> and create an account.
2. Go to <https://apps.twitter.com/> and log in
3. Click on "Create New App"
4. Give your app a name, description, and provide a website (you can use the course piazza site)
5. Agree to the terms and conditions
6. Go to the "Keys and Access Tokens" tab
7. Copy your Consumer Key (API) and Consumer Secret (API Secret) (DO NOT SHARE THESE!!)
8. Click on "Create my access token"
9. Copy your Access Token and Access Token Secret (DO NOT SHARE THESE!!)
10. The file getTweets.py contains skeleton code to filter tweets based on keywords. Download this file.
11. Update the file to include your Consumer Keys and Access Tokens
12. Fill in the missing code to add a field with the retrieval date and time and to store the tweet in the collection.

SUBMIT ONLY THE CODE YOU ENTERED where it says "YOUR CODE HERE"

13. Run the file and look at the content of the tweets

You can do this with the command

```
python getIceCreamTweets.py
```

or

```
python3 getIceCreamTweets.py
```

14. The program will continue to run until you kill it using Ctrl-C

Once you have your code running and have collected a number of tweets, write code to query the mongodb database.

Your code must generate an exact answer to the question provided. No human interpretation or intervention. You may write a script that produces the result, but you cannot, for example, write code that produces a document and then manually take a value from the document and plug it into more code.

1. (2 points) Add a field called "netId" with the value set to your netId to all of the records.
2. (2 points) How many tweets were longer than 140 characters?
3. (4 points) How many of the tweets contain the hashtag '#foodtruck'? Note that the tweet text can be in the text field or in the full_text field of the extended_tweet document.
4. (4 points) What are the top 5, non-whitespace only, user locations (exclude locations which contain only whitespace or tabs)? Display the location value and the counts. Subsort by location, if there are ties. You might find this webpage helpful: https://en.wikipedia.org/wiki/Regular_expression for determining a regular expression for Space and tab.
5. (2 points) How many of the tweets were retweeted?
6. (4 points) What how many followers does the tweeter with the most followers have?

7. (4 points) Use mongoexport to export your tweets to a json file. Name the file icecream<yourNetIdHere>.json. List your export statement here.

A skeleton is:

```
mongoexport --db ricedb --collection icecream
--out icecream<yourNetIdHere>.json --username ricedb
--password "myPwd"
```

3 Short Answers

1. (2 points) List 1 advantage and 1 disadvantage of embedding documents within other documents
2. (2 points) The tweets contain a number of fields with 0 or null values. What is an advantage of including these fields?
3. (2 points) What is a disadvantage?

4 What to submit

Submit a zip file containing:

1. A .js file that contains your mongodb javascript executable code. Basically, we want to be able to cut & paste your code into the Mongodb client and have it run, or run it at the command line. This means that any comments or text answers in this file should be in Mongodb comments. A line is considered a comment if it starts with "//".
2. Your modified getIceCreamTweets.py file (you can take out the access codes)
3. (5 points) A json file containing the tweets you collected
4. A text .txt file with your answers